

Processor Pipeline

Instructor: Nima Honarmand






Processor Datapath

The Generic Instruction Cycle

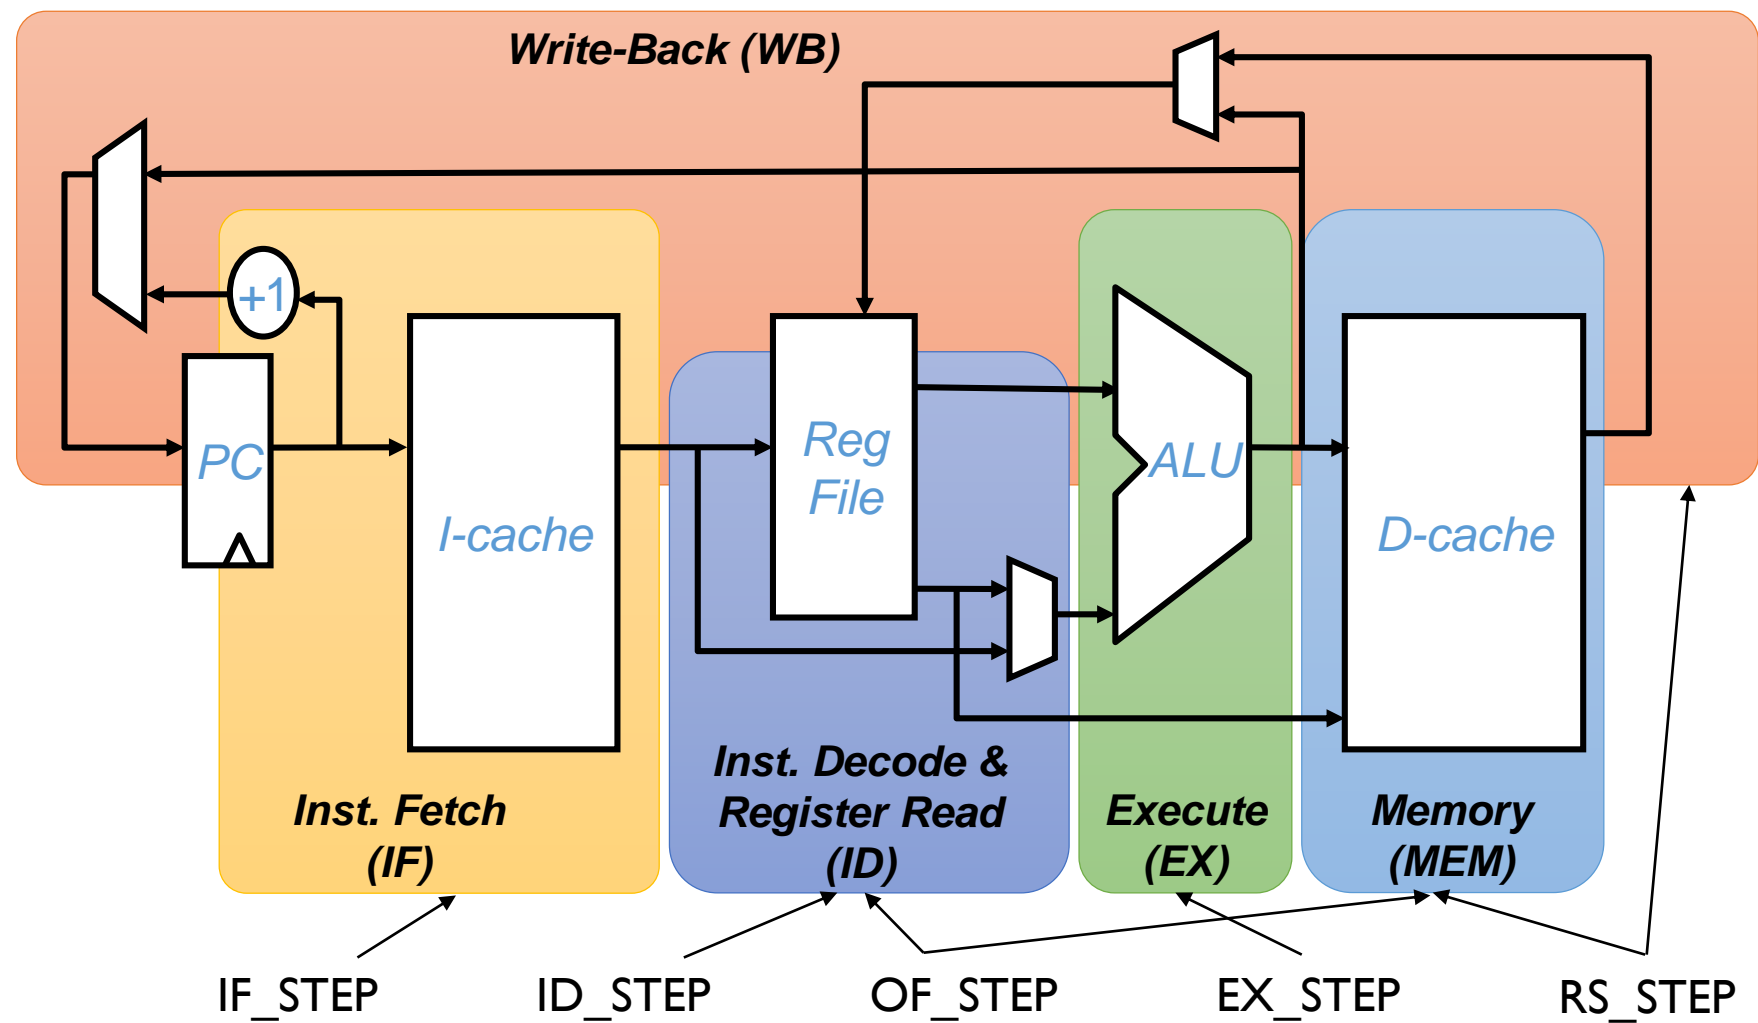
- Steps in processing an instruction:
 - Instruction Fetch (**IF_STEP**)
 - Instruction Decode (**ID_STEP**)
 - Operand Fetch (**OF_STEP**)
 - Might be from registers or memory
 - Execute (**EX_STEP**)
 - Perform computation on the operands
 - Result Store or Write Back (**RS_STEP**)
 - Write the execution results back to registers or memory
- ISA determines what needs to be done in each step for each instruction
- μ Arch determines how HW implements the steps

Example: MIPS Instruction Set

- All instructions are 32 bits



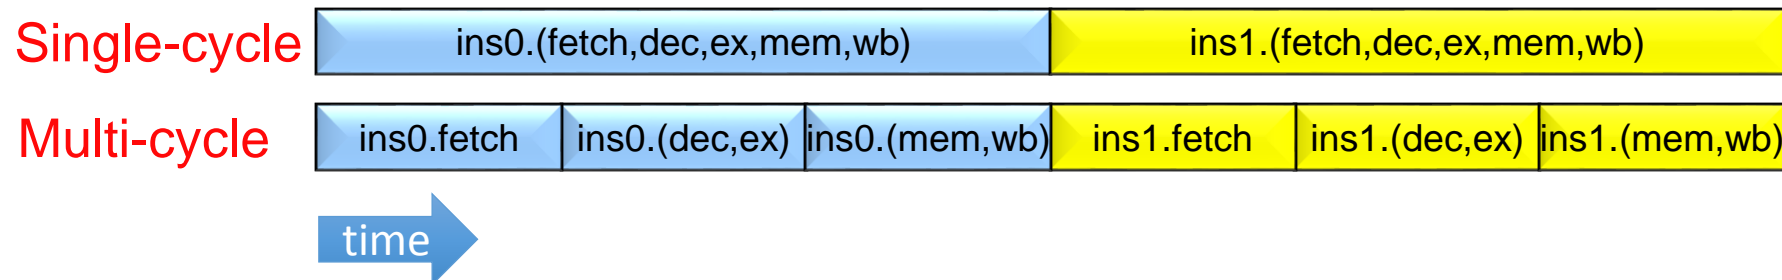
A Simple MIPS Datapath



Datapath vs. Control Logic

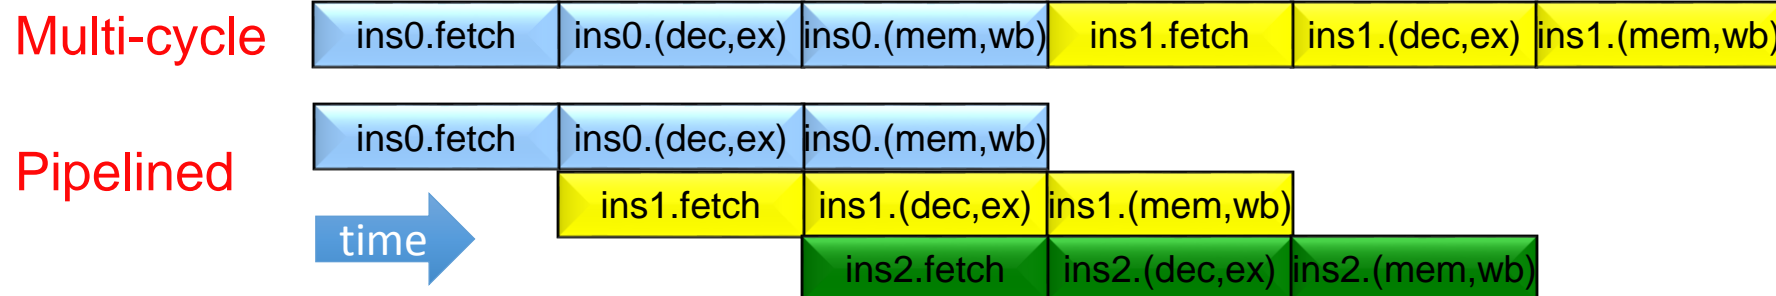
- **Datapath** is the collection of HW components and their connection in a processor
 - Determines the *static* structure of processor
- **Control logic** determines the dynamic flow of data between the components
 - E.g., the control lines of MUXes and ALU in last slide
 - Is a function of?
 - Instruction words
 - State of the processor
 - Execution results at each stage

Single-Instruction Datapath



- Process one instruction at a time
- Single-cycle control: hardwired
 - Low CPI (1)
 - Long clock period (to accommodate slowest instruction)
- Multi-cycle control: typically micro-programmed
 - Short clock period
 - High CPI
- Can we have both low CPI and short clock period?
 - Not if datapath executes only one instruction at a time
 - No good way to make a single instruction go faster

Pipelined Datapath





- Start with multi-cycle design
- When insn0 goes from stage 1 to stage 2
... insn1 starts stage 1
- Each instruction passes through all stages
... but instructions enter and leave at faster *rate*

Style	Ideal CPI	Cycle Time (1/freq)
Single-cycle	1	Long
Multi-cycle	> 1	Short
Pipelined	1	Short

Pipeline can have as many insns *in flight* as there are stages

Pipeline Illustrated



Pipeline Latency = n Gate Delay + $(p-1)$ register delays
 p : # of stages

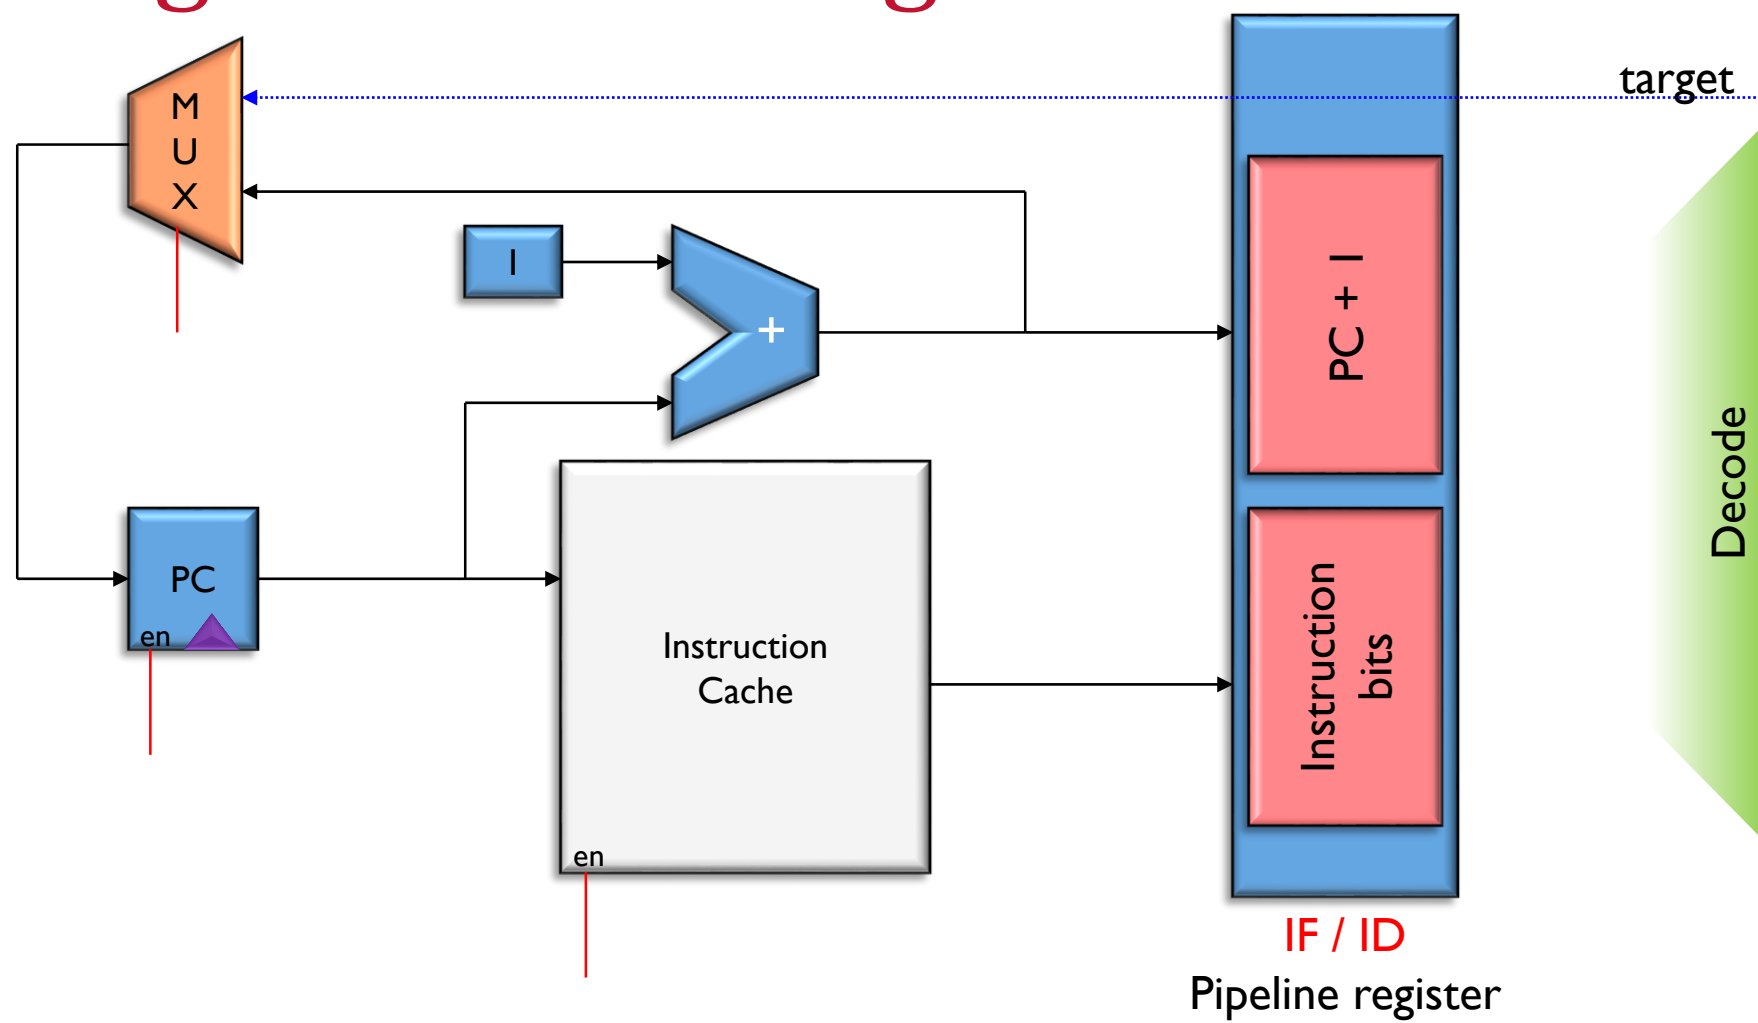
Improves throughput at the expense of latency

5-Stage MIPS Pipeline

Stage 1: Fetch

- Fetch an instruction from instruction cache every cycle
 - Use PC to index instruction cache
 - Increment PC (assume no branches for now)
- Write state to the pipeline register (IF/ID)
 - The next stage will read this pipeline register

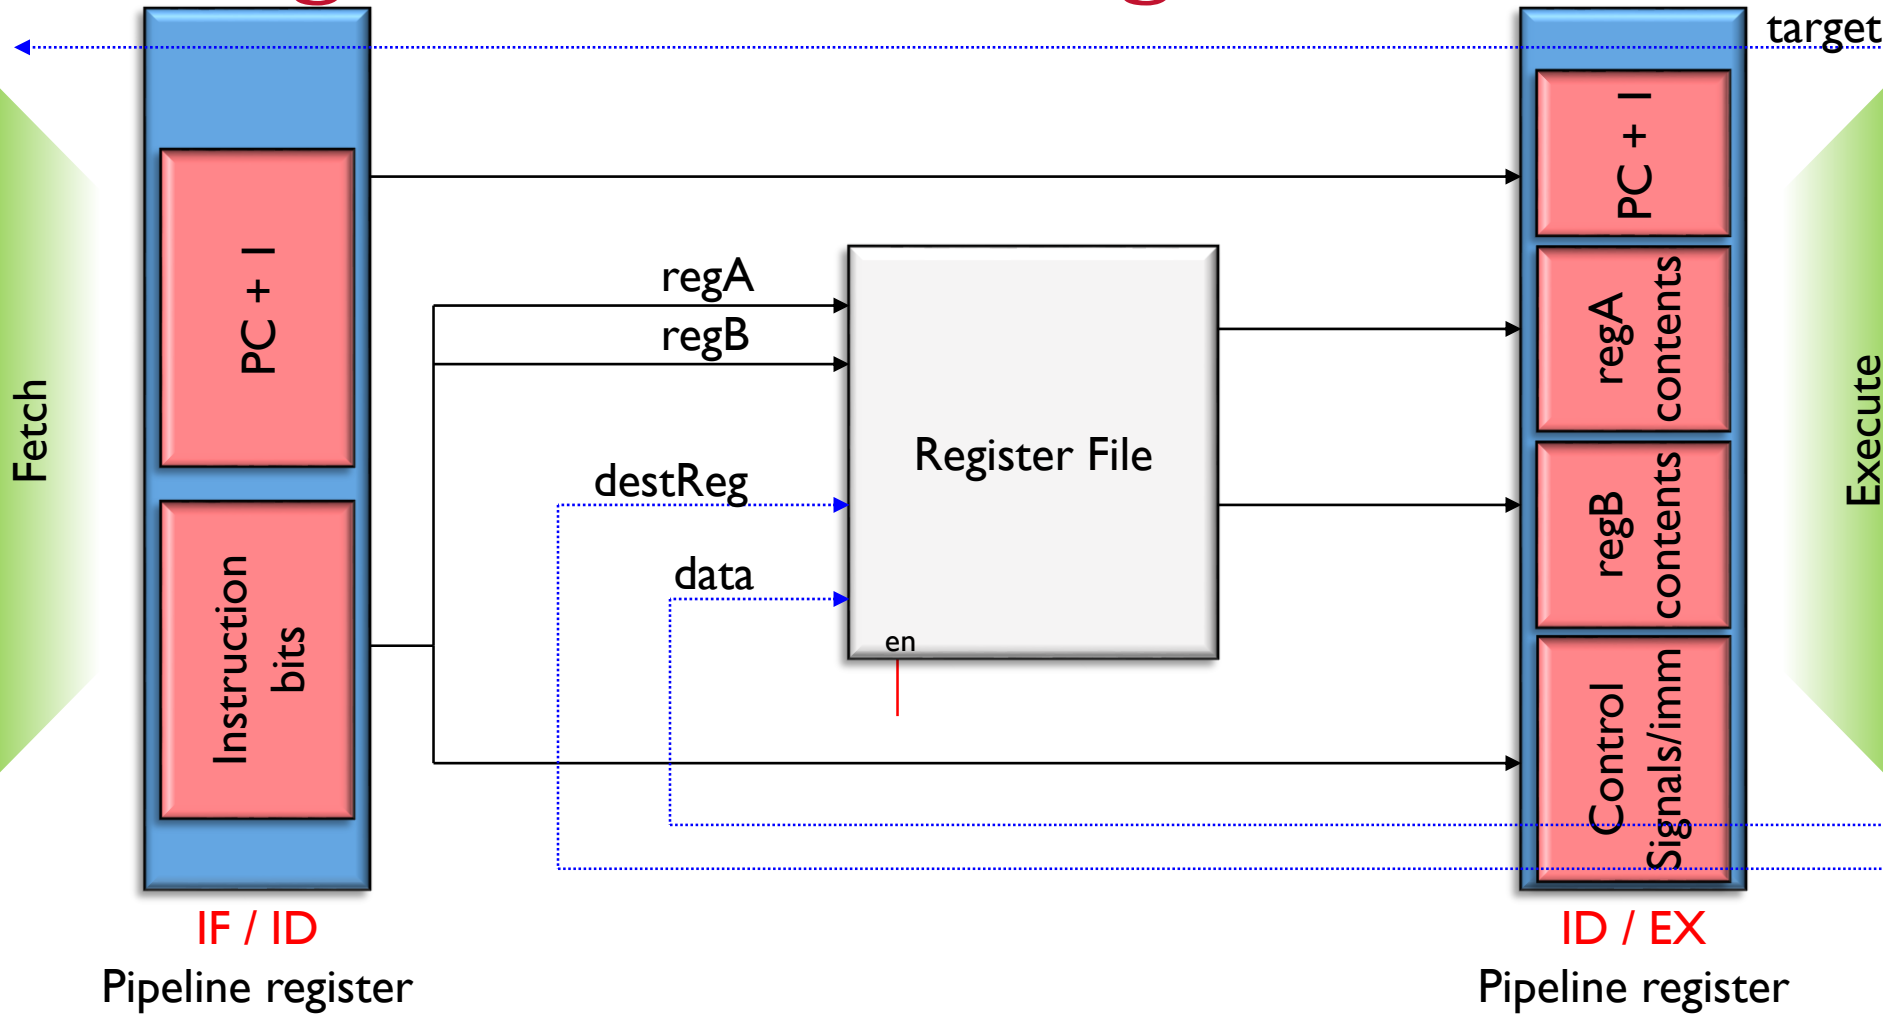
Stage 1: Fetch Diagram



Stage 2: Decode

- Decodes opcode bits
 - Set up Control signals for later stages
- Read input operands from register file
 - Specified by decoded instruction bits
- Write state to the pipeline register (ID/EX)
 - Opcode
 - Register contents, immediate operand
 - PC+1 (even though decode didn't use it)
 - Control signals (from insn) for opcode and destReg

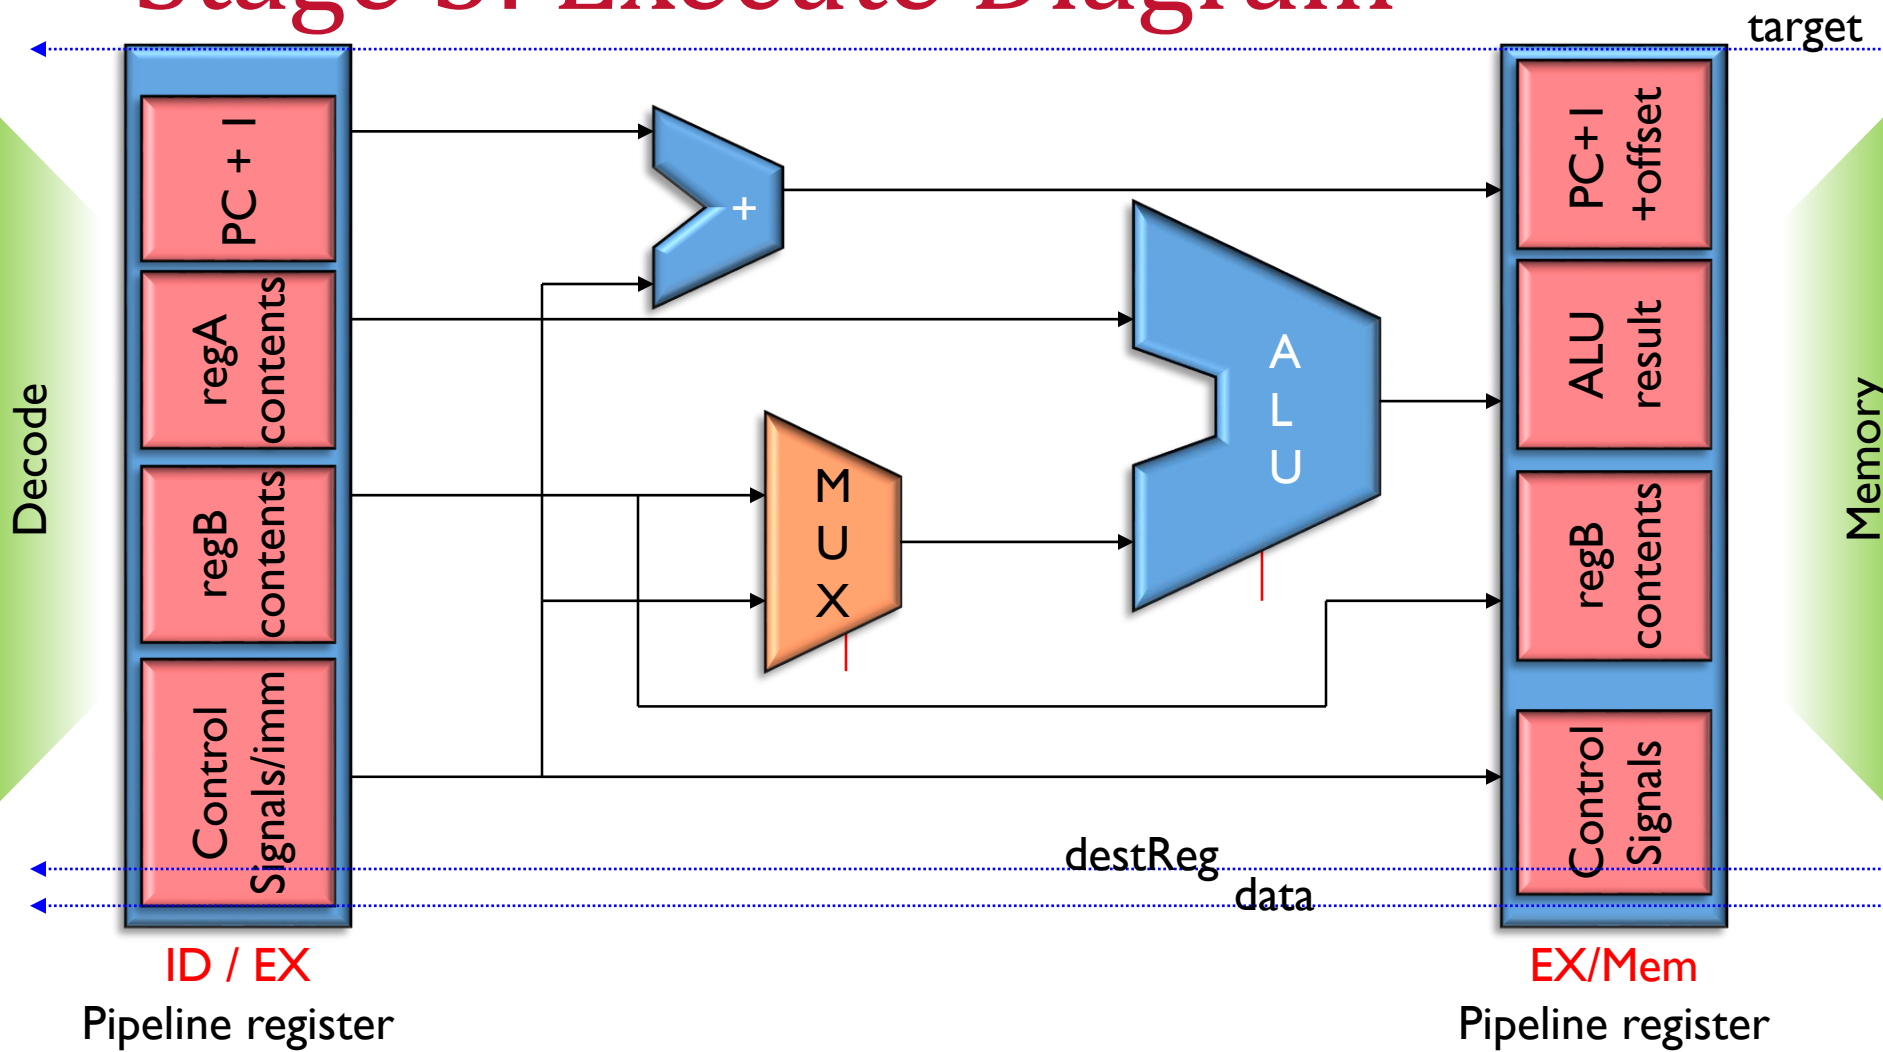
Stage 2: Decode Diagram



Stage 3: Execute

- Perform ALU operations
 - Calculate result of instruction
 - Control signals select operation
 - Contents of regA used as one input
 - Either regB or constant offset (imm from insn) used as second input
 - Calculate PC-relative branch target
 - $PC+1+(\text{constant offset})$
- Write state to the pipeline register (EX/Mem)
 - ALU result, contents of regB, and $PC+1+\text{offset}$
 - Control signals (from insn) for opcode and destReg

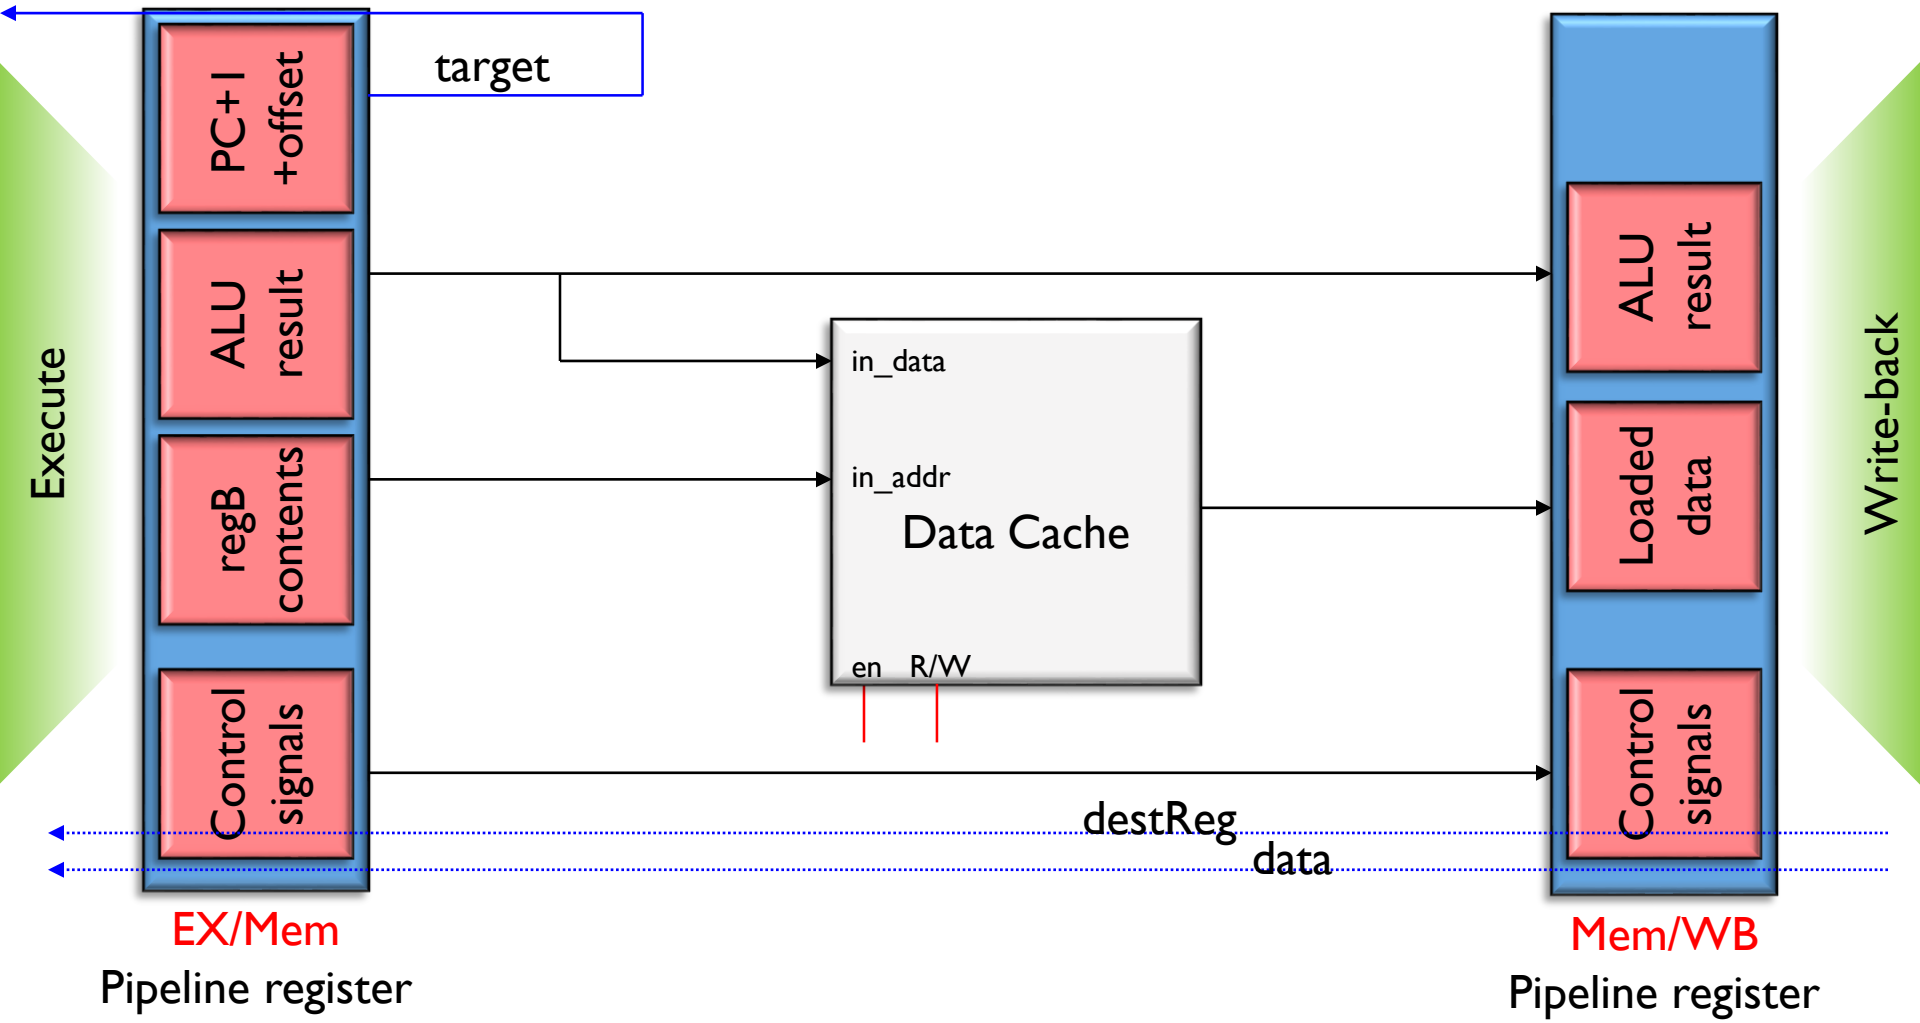
Stage 3: Execute Diagram



Stage 4: Memory

- Perform data cache access
 - ALU result contains address for LD or ST
 - Opcode bits control R/W and enable signals
- Write state to the pipeline register (Mem/WB)
 - ALU result and Loaded data
 - Control signals (from insn) for opcode and destReg

Stage 4: Memory Diagram

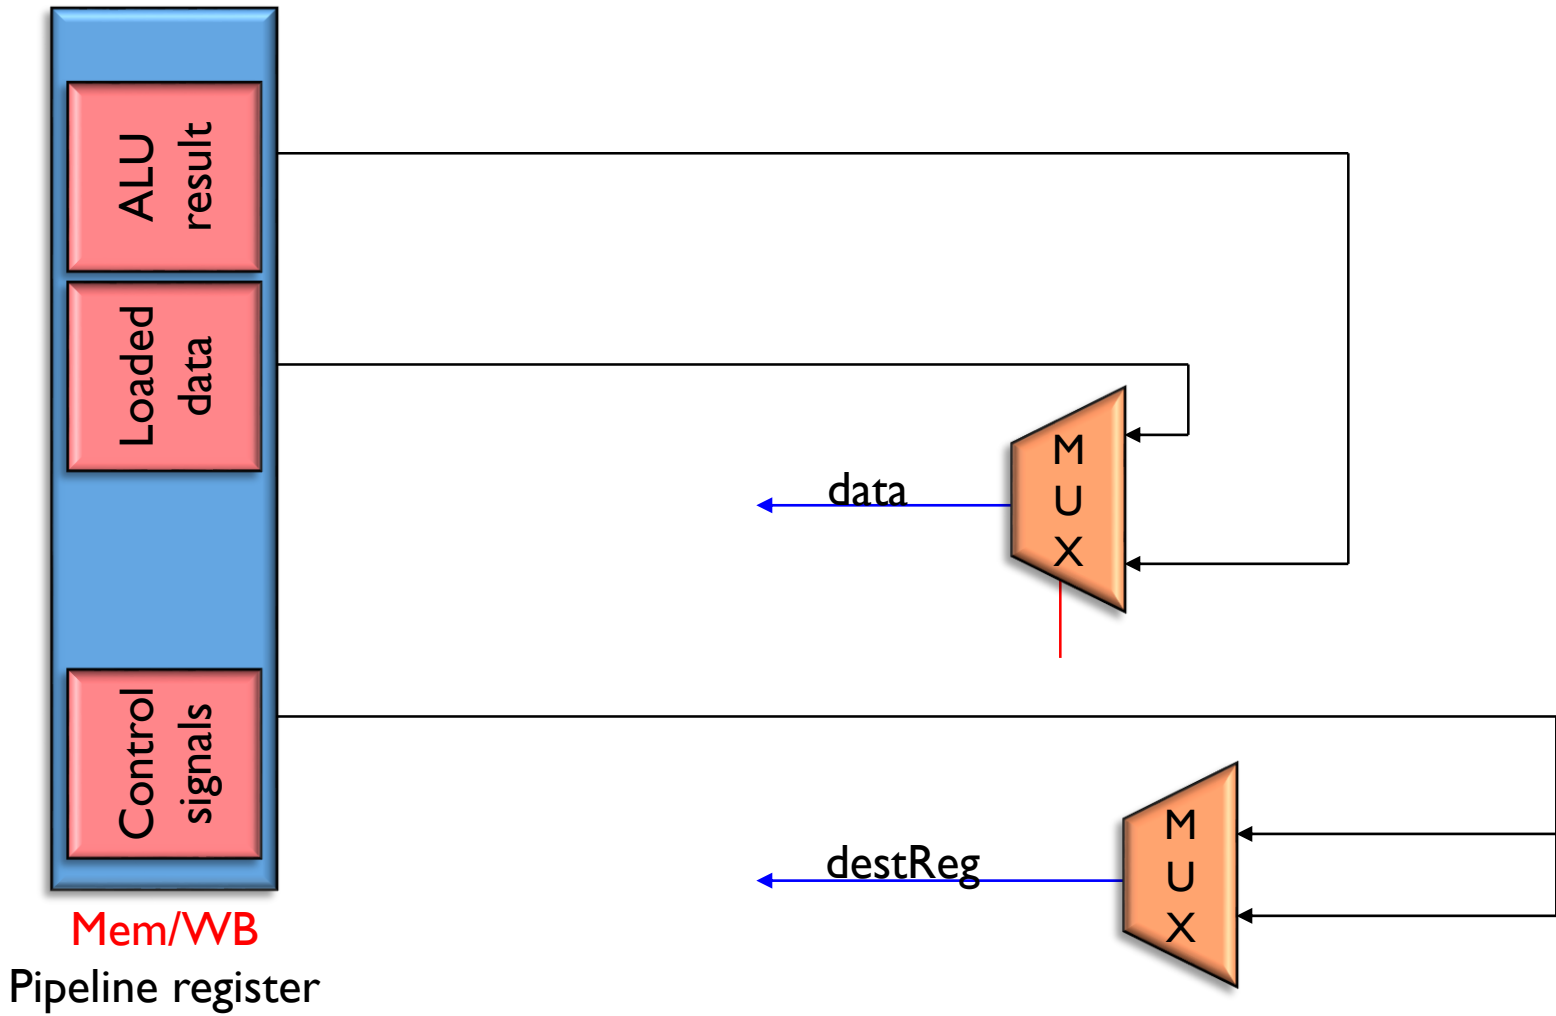


Stage 5: Write-back

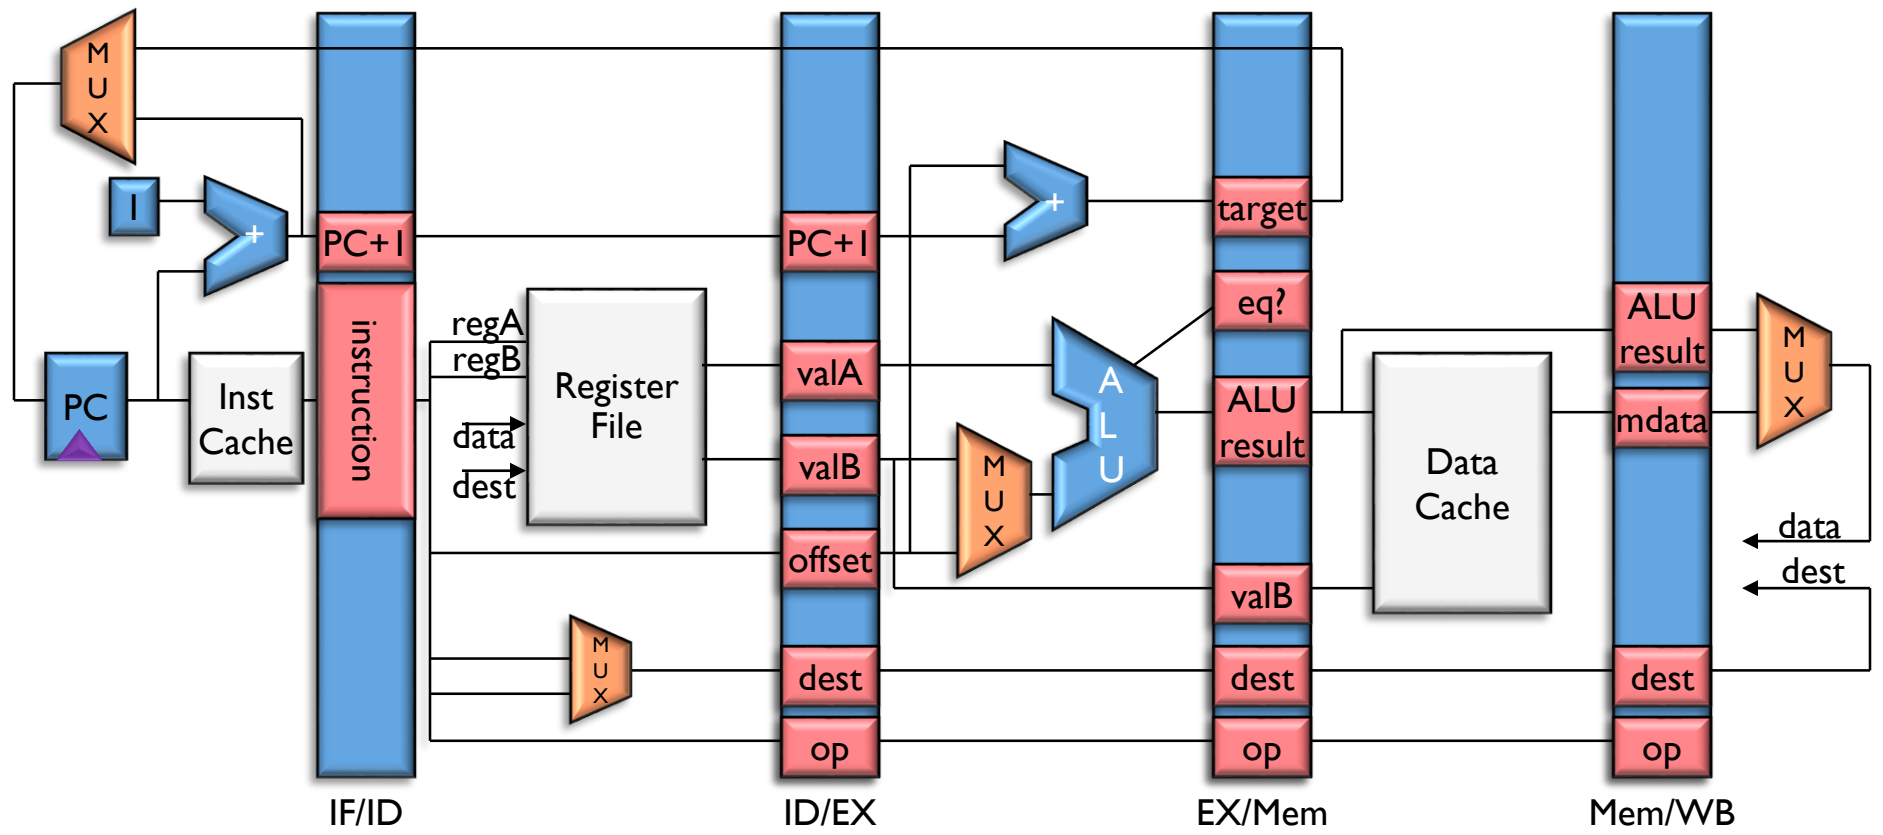
- Writing result to register file (if required)
 - Write Loaded data to destReg for LD
 - Write ALU result to destReg for ALU insn
 - Opcode bits control register write enable signal

Stage 5: Write-back Diagram

Memory



Putting It All Together



Issues With Pipelining

Pipelining Idealism


- Uniform Sub-operations
 - Operation can partitioned into uniform-latency sub-ops
- Repetition of Identical Operations
 - Same ops performed on many different inputs
- Independent Operations
 - All ops are mutually independent

Pipeline Realism

- Uniform Sub-operations ... NOT!
 - Balance pipeline stages
 - Stage quantization to yield balanced stages
 - Minimize internal fragmentation (left-over time near end of cycle)
- Repetition of Identical Operations ... NOT!
 - Unifying instruction types
 - Coalescing instruction types into one “multi-function” pipe
 - Minimize external fragmentation (idle stages to match length)
- Independent Operations ... NOT!
 - Resolve data and resource hazards
 - Inter-instruction dependency detection and resolution

Pipelining is expensive


The Generic Instruction Pipeline



Balancing Pipeline Stages



$T_{IF} = 6$ units



$T_{ID} = 2$ units



$T_{OF} = 9$ units



$T_{EX} = 5$ units



$T_{OS} = 9$ units

Without pipelining

$$T_{cyc} \approx T_{IF} + T_{ID} + T_{OF} + T_{EX} + T_{OS} \\ = 31$$

Pipelined

$$T_{cyc} \approx \max\{T_{IF}, T_{ID}, T_{OF}, T_{EX}, T_{OS}\} \\ = 9$$

$$\text{Speedup} = 31 / 9 = 3.44$$

Can we do better?

Balancing Pipeline Stages (1/2)

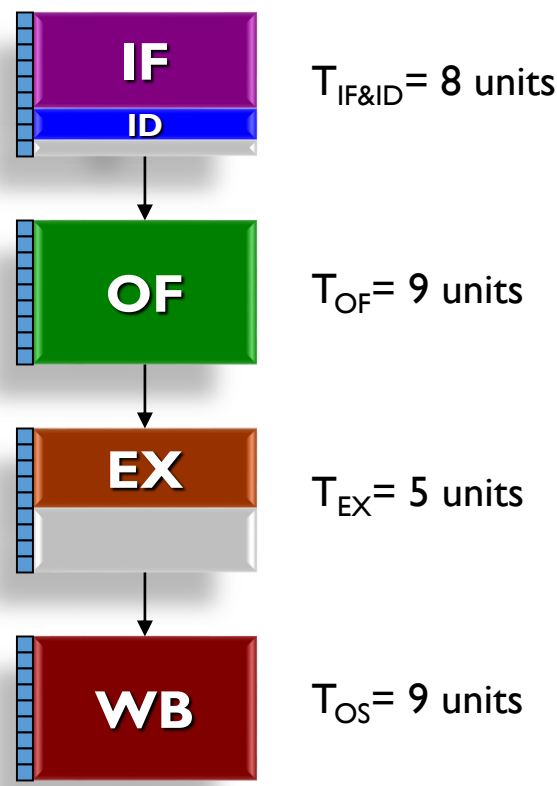
- Two methods for stage quantization
 - Divide sub-ops into smaller pieces
 - Merge multiple sub-ops into one
- Recent/Current trends
 - Deeper pipelines (more and more stages)
 - Pipelining of memory accesses
 - Multiple different pipelines/sub-pipelines

Balancing Pipeline Stages (2/2)


Coarser-Grained Machine Cycle:
4 machine cyc / instruction

Finer-Grained Machine Cycle:
11 machine cyc /instruction

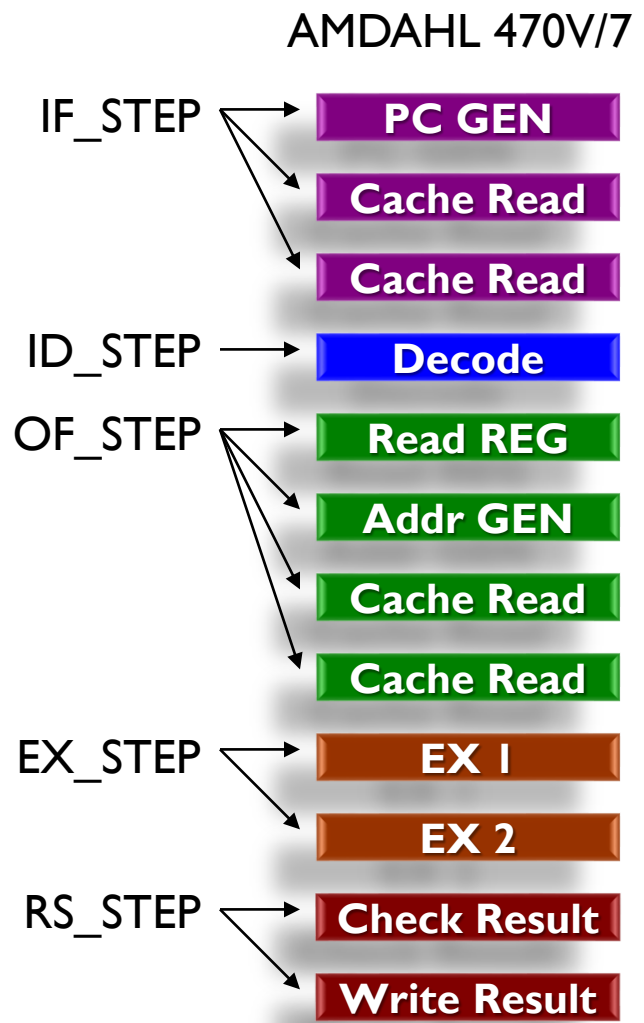

stages = 4
 $T_{cyc} = 9$ units



stages = 11
 $T_{cyc} = 3$ units



Pipeline Examples




Instruction Dependencies (1/2)

- Data Dependence
 - Read-After-Write (RAW) (the only true dependence)
 - Read must wait until earlier write finishes
 - Anti-Dependence (WAR)
 - Write must wait until earlier read finishes (avoid clobbering)
 - Output Dependence (WAW)
 - Earlier write can't overwrite later write
- Control Dependence (a.k.a. Procedural Dependence)
 - Branch condition must execute before branch target
 - Instructions after branch cannot run before branch

Instruction Dependencies (1/2)

From
Quicksort: # for (; (j < high) && (array[j] < array[low]); ++j);



Hardware Dependency Analysis

- Processor must handle
 - Register Data Dependencies (same register)
 - RAW, WAW, WAR
 - Memory Data Dependencies (same address)
 - RAW, WAW, WAR
 - Control Dependencies

Pipeline Terminology

- Pipeline Hazards
 - Potential violations of program dependencies
 - Due to multiple in-flight instructions
 - Must ensure program dependencies are not violated
- Hazard Resolution
 - Static method: compiler guarantees correctness
 - By inserting No-Ops or independent insns between dependent insns
 - Dynamic method: hardware checks at runtime
 - Two basic techniques: **Stall** (costs perf.), **Forward** (costs hw)
- Pipeline Interlock
 - Hardware mechanism for dynamic hazard resolution
 - Must detect and enforce dependencies at runtime


The diagram illustrates a 5-stage pipeline (IF, ID, RD, ALU, MEM, WB) over six time steps (t_0 to t_5). Five instructions ($Inst_j$ to $Inst_{j+4}$) are shown executing in parallel. A horizontal blue arrow at the top points right, and a vertical blue arrow on the left points down.

	t_0	t_1	t_2	t_3	t_4	t_5
$Inst_j$	IF	ID	RD	ALU	MEM	WB
$Inst_{j+1}$		IF	ID	RD	ALU	MEM
$Inst_{j+2}$			IF	ID	RD	ALU
$Inst_{j+3}$				IF	ID	RD
$Inst_{j+4}$					IF	ID

Data Hazards

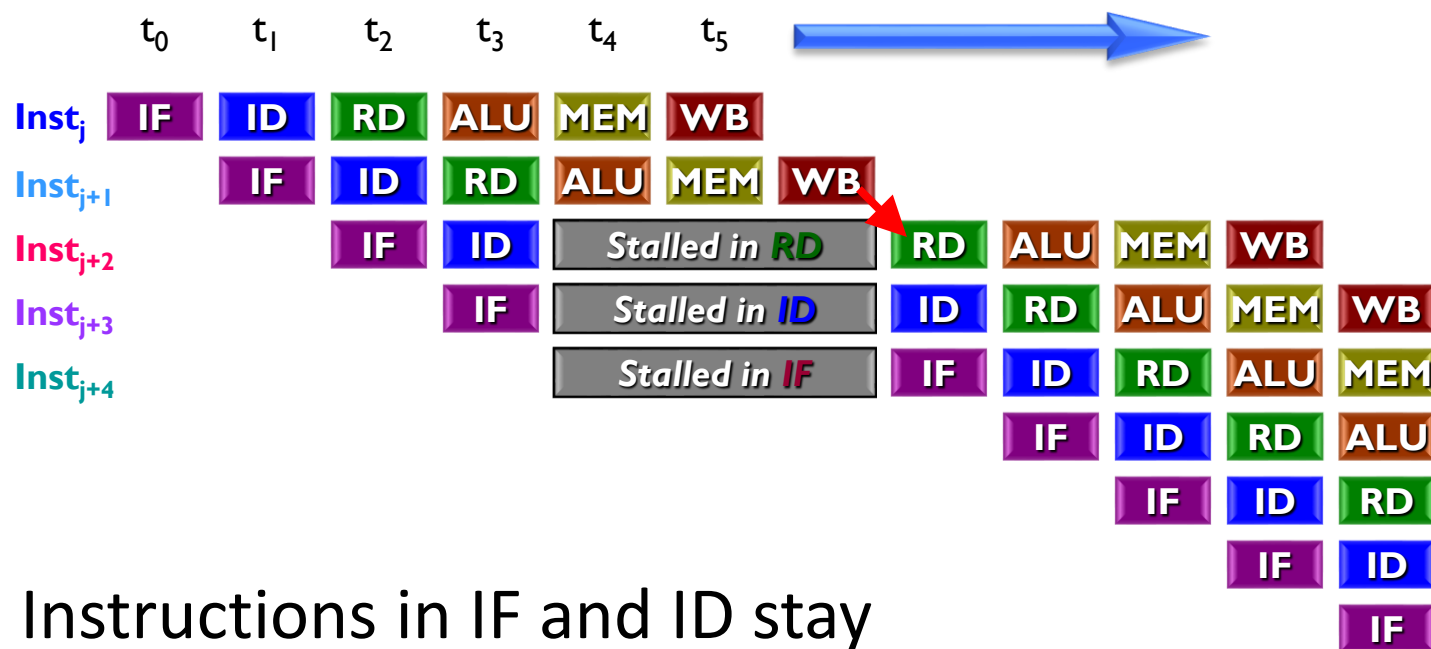
- Necessary conditions:
 - WAR: write stage earlier than read stage
 - Is this possible in IF-ID-RD-EX-MEM-WB?
 - WAW: write stage earlier than write stage
 - Is this possible in IF-ID-RD-EX-MEM-WB?
 - RAW: read stage earlier than write stage
 - Is this possible in IF-ID-RD-EX-MEM-WB?
- If conditions not met, no need to resolve
- Check for both register and memory

Pipeline: Data Hazard




- Only RAW in our case
- How to detect?
 - Compare read register specifiers for newer instructions with write register specifiers for older instructions

Option 1: Stall on Data Hazard



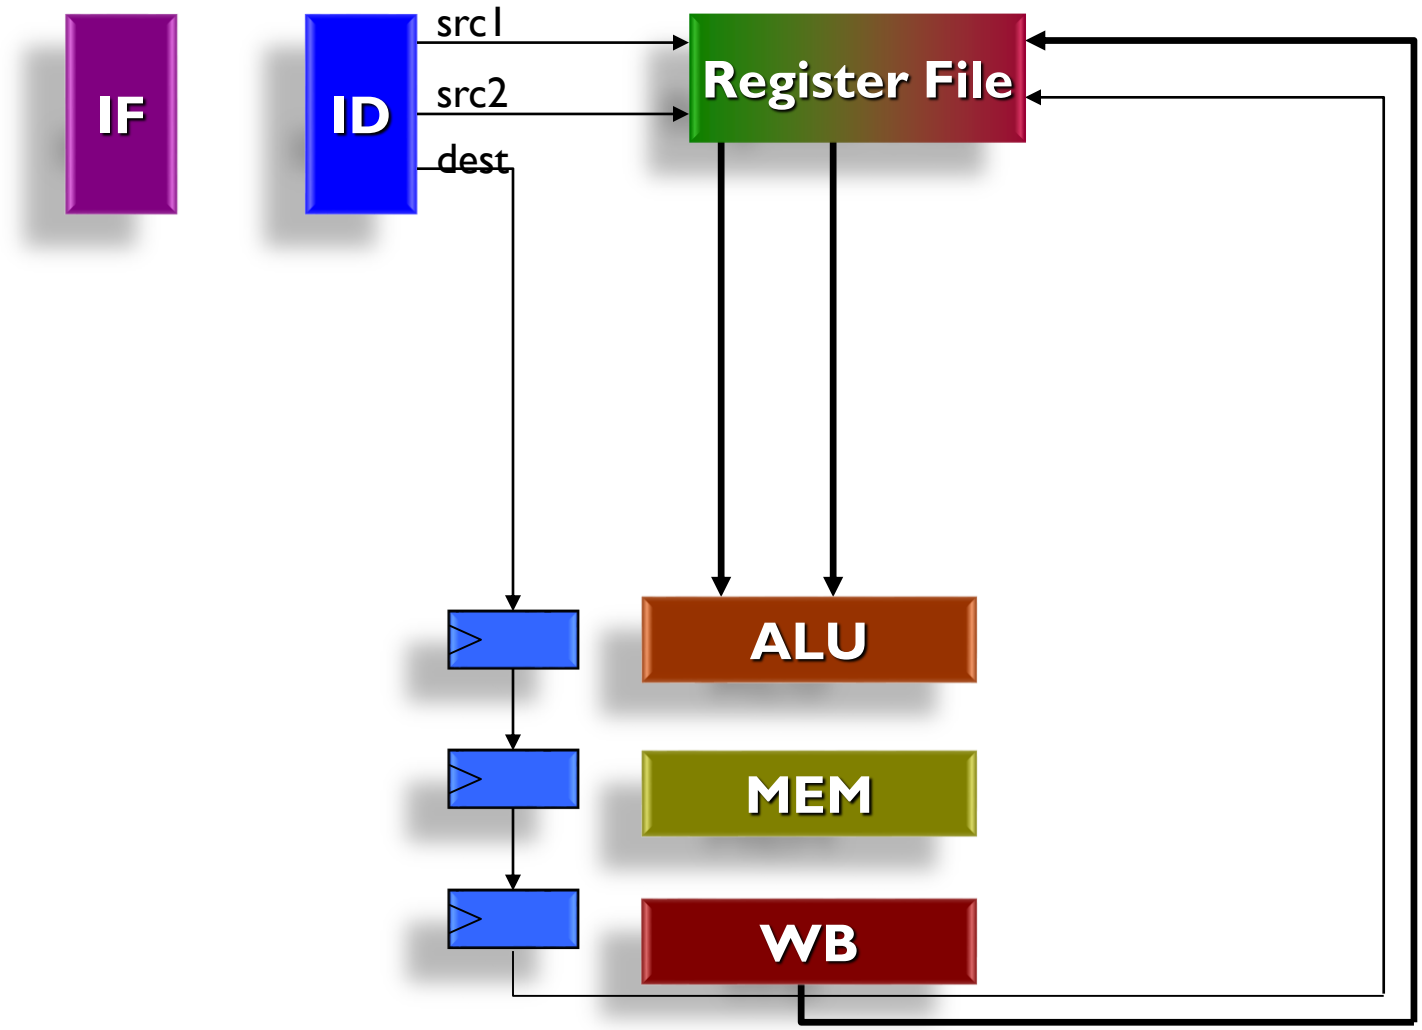
- Instructions in IF and ID stay
- IF/ID pipeline latch not updated
- Send no-op down pipeline (called a bubble)

Option 2: Forwarding Paths (1/3)

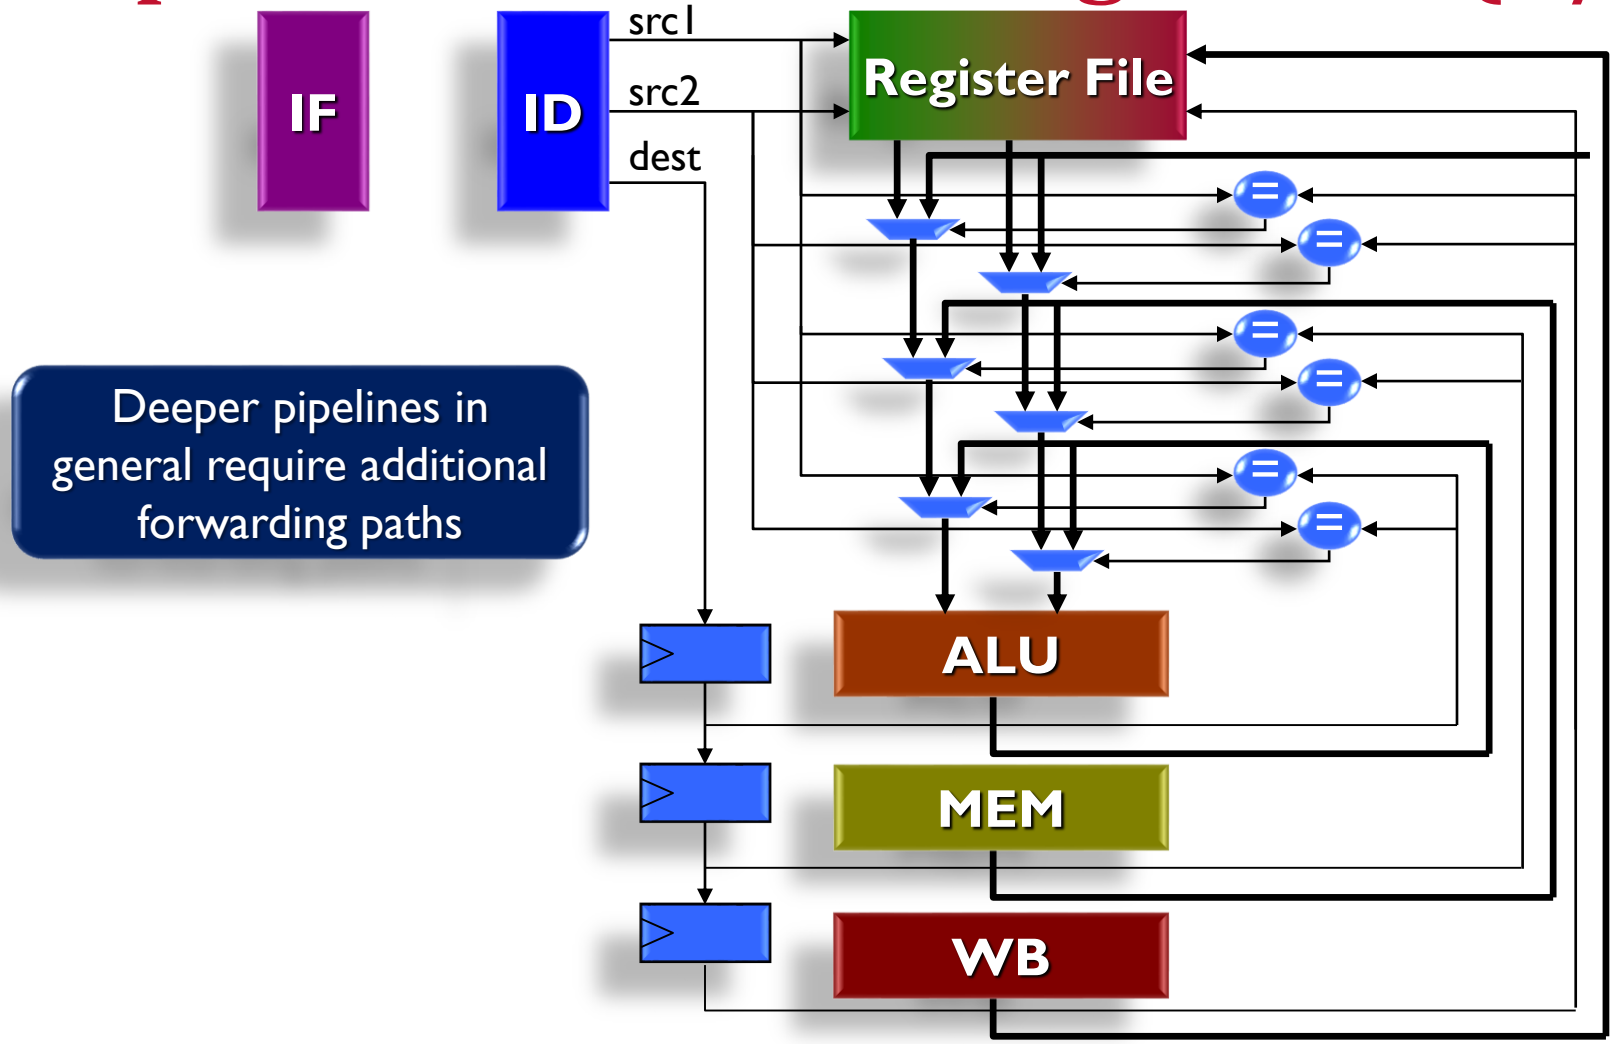


Requires stalling even with forwarding paths


Option 2: Forwarding Paths (2/3)



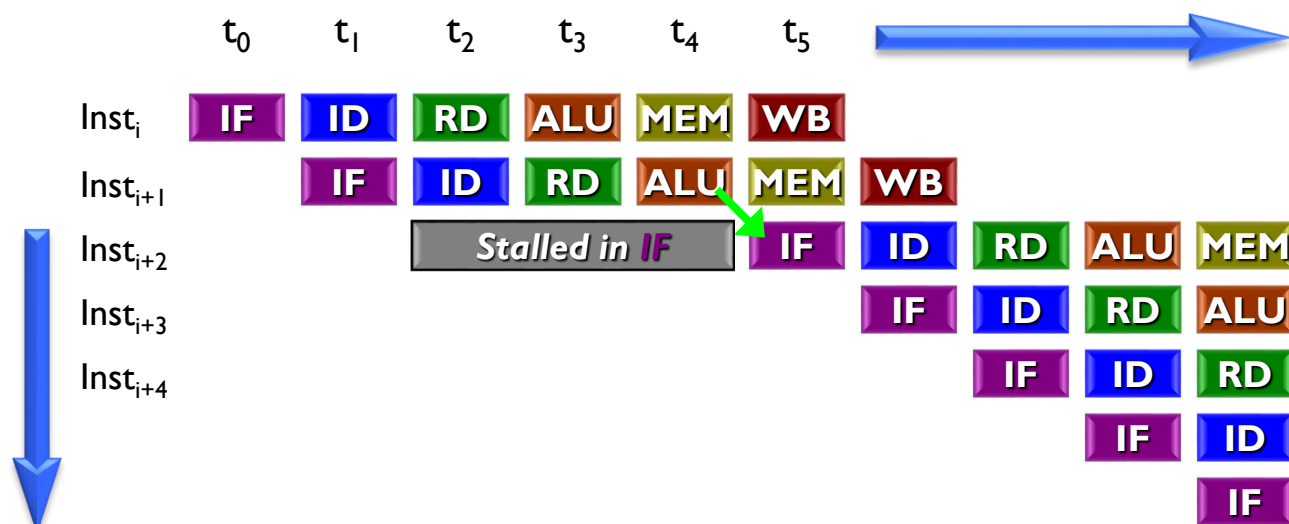
Option 2: Forwarding Paths (3/3)



Pipeline: Control Hazard




Option 1: Stall on Control Hazard



- Stop fetching until branch outcome is known
 - Send no-ops down the pipe
- Easy to implement
- Performs poorly
 - On out of 6 instructions are branches
 - Each branch takes 3 cycles
 - $CPI = 1 + 3 \times 1/6 = 1.5$ (lower bound)

Option 2: Prediction for Control Hazards



- Predict branch not taken
- Send sequential instructions down pipeline
- Must stop memory and RF writes
- Kill instructions later if incorrect
- Fetch from branch target

Option 3: Delay Slots for Control Hazards

- Another option: delayed branches
 - # of delay slots (ds) : stages between IF and where the branch is resolved
 - 3 in our example
 - Always execute following ds instructions
 - Put useful instruction there, otherwise no-op
- Losing popularity
 - Just a stopgap (one cycle, one instruction)
 - Superscalar processors (later)
 - Delay slot just gets in the way (special case)

Legacy from old RISC ISAs

Superscalar Pipelines


Instruction-Level Parallelism Beyond Simple Pipelines

Going Beyond Scalar

- Scalar pipeline limited to $\text{CPI} \geq 1.0$
 - Can never run more than 1 insn per cycle
- “Superscalar” can achieve $\text{CPI} \leq 1.0$ (i.e., $\text{IPC} \geq 1.0$)
 - Superscalar means executing multiple insns in parallel


Architectures for Instruction Parallelism

- Scalar pipeline (baseline)
 - Instruction/overlap parallelism = D
 - Operation Latency = 1
 - Peak IPC = 1.0

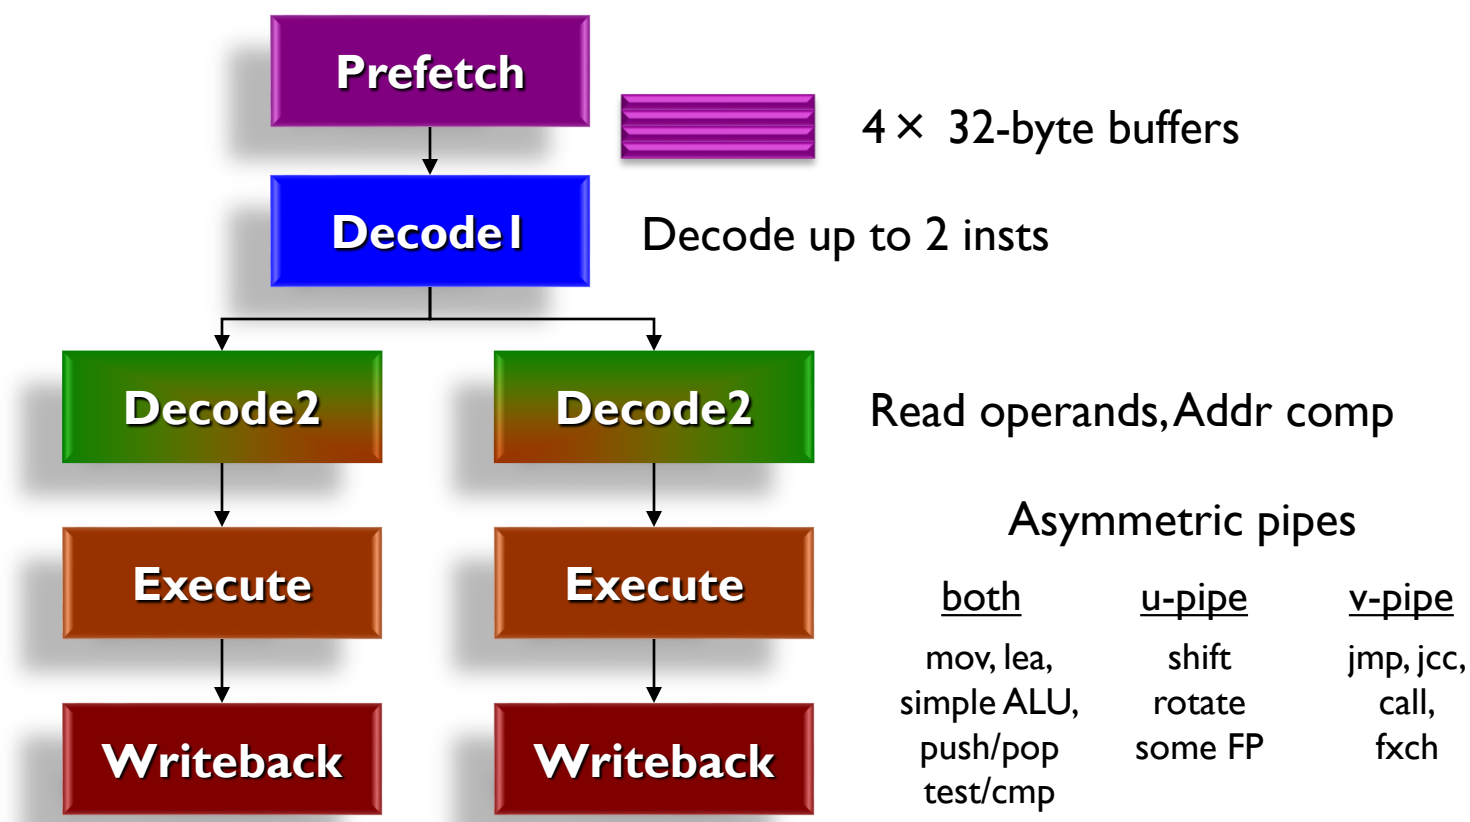


Superscalar Machine

- Superscalar (pipelined) Execution
 - Instruction parallelism = $D \times N$
 - Operation Latency = 1
 - Peak IPC = N per cycle



Superscalar Example: Pentium



Pentium Hazards & Stalls

- “Pairing Rules” (when can’t two insns exec?)
 - Read/flow dependence
 - `mov eax, 8`
 - `mov [ebp], eax`
 - Output dependence
 - `mov eax, 8`
 - `mov eax, [ebp]`
 - Partial register stalls
 - `mov al, 1`
 - `mov ah, 0`
 - Function unit rules
 - Some instructions can never be paired
 - `MUL`, `DIV`, `PUSHA`, `MOVS`, some FP

Limitations of In-Order Pipelines

- If the machine parallelism is increased
 - ... dependencies reduce performance
 - CPI of in-order pipelines degrades sharply
 - As N approaches avg. distance between dependent instructions
 - Forwarding is no longer effective
 - Must stall often

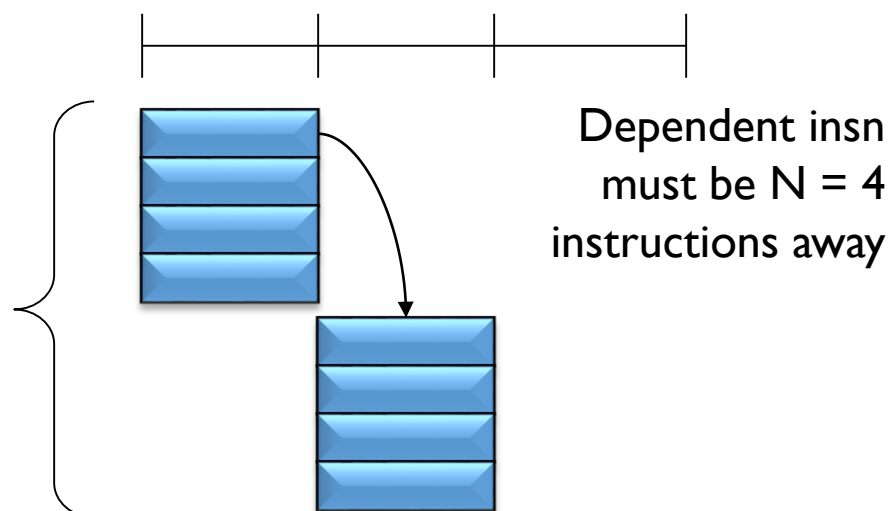
In-order pipelines are rarely full

The In-Order N-Instruction Limit

- On average, parent-child separation is about 5 insn
– (Franklin and Sohi '92)

Ex. Superscalar degree $N = 4$

Any dependency
between these
instructions will
cause a stall



Average of 5 means there are many
cases when the separation is < 4 ...
each of these limits parallelism

Reasonable in-order superscalar is effectively $N=2$